

Definitions:

Pearson's correlation coefficient:

The covariance of two variables divided by the product of their standard deviations.

For a population:

$$\rho_{x,y} = \frac{\text{Cov}(X, Y)}{\sigma_X \sigma_Y}$$

where

$$\text{Cov}(X, Y) = E[(X - E[X])(Y - E[Y])]$$

For a sample:

It's often referred to as the sample correlation coefficient, commonly abbreviated to just "r"

$$r = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}$$

(Above: the sample covariance divided by the product of the sample standard deviations)

which can be manipulated to get:

$$r = r_{xy} = \frac{1}{n-1} \sum_{i=1}^n \left(\frac{x_i - \bar{x}}{s_x} \right) \left(\frac{y_i - \bar{y}}{s_y} \right)$$

The correlation coefficient always takes a value between -1 and 1, with 1 or -1 indicating perfect correlation (all points would lie along a straight line in this case).

- A positive correlation indicates a positive association between the variables (increasing values in one variable correspond to increasing values in the other variable).
- A negative correlation indicates a negative association between the variables (increasing values in one variable correspond to decreasing values in the other variable).
- A correlation value close to 0 indicates no association between the variables.

The square of the correlation coefficient, R^2 , is a useful value in linear regression. This value represents the fraction of the variation in one variable that may be explained by the other variable. Thus, if a correlation of $r = 0.8$ is observed between two variables (say, height and weight, for example), then a linear regression model attempting to explain either variable in terms of the other variable will account for 64% ($r^2 = 0.8^2 = .64$) of the variability in the data.

The correlation coefficient also relates directly to the regression line $Y = a + bX$ for any two variables, where $b = r \frac{s_x}{s_y}$

I found this info here:

<http://www.stat.yale.edu/Courses/1997-98/101/correl.htm>

and on the wikipedia page.

Regression Analysis definition

A statistical technique for modeling and investigating the relationship between variables.

The basic model is:

$$y = \beta_0 + \beta_1 x + \epsilon$$

The **response variable**, y , is the variable you're analyzing to see how much it's influenced by the other variable(s).

The **regressor variable(s)**, x , is (are) the variable(s) you're estimating regression coefficients for in order to predict future response variables.

The **regression coefficients**, β_0, β_1, \dots are the coefficients for each regressor variable (and a slope, usually) that best minimize the random error (ϵ) for the model.

The **random error** term, ϵ , is the random variable that accounts for the failure of the model to fit the data exactly. For example, for a particular (x_i, y_i) , the ϵ_i is

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i$$

where

$$\epsilon \sim N(0, \sigma^2)$$

The expected values of each of these quantities are:

$$E[y | x] = \mu_{y|x} = E[\beta_0 + \beta_1 x + \epsilon] = E[\beta_0] + E[\beta_1 x] + E[\epsilon] = \beta_0 + \beta_1 x + 0$$

$$V[y | x] = \sigma_{y|x}^2 = V[\beta_0 + \beta_1 x + \epsilon] = V[\beta_0] + V[\beta_1 x] + V[\epsilon] = 0 + 0 + \sigma^2 = \sigma^2$$

What're the 3 key assumptions?

- a. Uncorrelated Errors (what does this mean specifically?)
- b. Constant Variance (**between what?**)
- c. and one other...

Constructing a Regression model:

The β 's must all be estimated.

For a sample regression model:

$$y_i = \beta_0 + \beta_1 x_{i,1} + \dots + \beta_{i,k} x_{i,k} + \epsilon_i \text{ for } i = 0, 1, 2 \dots n$$

Least squares estimation seeks to minimize the sum of the squares of the differences between the observed responses (the y_i 's) and the straight line.

$$S(\beta_0, \beta_1, \dots) = \sum \epsilon_i^2 = \sum [y - (\beta_0 + \beta_1 x_{i,1} + \dots)]^2$$

When you take the partial derivative of each β , you get $k + 1$ equations. Since you have $k + 1$ unknowns, you can do some linear algebra to solve for each β . In the case of simple linear regression:

$$\begin{aligned} \hat{\beta}_0 &= \bar{y} - \hat{\beta}_1 \bar{x} \\ \hat{\beta}_1 &= \frac{\sum_{i=1}^n y_i x_i - \frac{(\sum_{i=1}^n y_i)(\sum_{i=1}^n x_i)}{n}}{\sum_{i=1}^n x_i^2 - \frac{(\sum_{i=1}^n x_i)^2}{n}} \end{aligned}$$

Put another way:

$$\begin{aligned} S_{xx} &= \sum_{i=1}^n (x_i - \bar{x})^2 \\ S_{xy} &= \sum_{i=1}^n (x_i - \bar{x}) y_i \\ \hat{\beta}_1 &= \frac{S_{xy}}{S_{xx}} \end{aligned}$$

Residuals

The **residuals** of a linear regression model are the errors for each sample which will later be used to determine the adequacy of the model.

$$\epsilon_i = y_i - \hat{y}_i$$

Some properties of the Least Squares Estimators (2.2.2)

The ordinary least-squares (OLS) estimator of the slope ($\hat{\beta}_1$) is a linear combination of the observations, y_i :

$$\hat{\beta}_1 = \frac{S_{xy}}{S_{xx}} = \sum_{i=1}^n c_i y_i$$

where

$$c_i = \frac{(x_i - \bar{x})}{S_{xx}}, \quad \sum_{i=1}^n c_i = 0, \quad \sum_{i=1}^n c_i^2 = \frac{1}{S_{xx}}, \quad \sum_{i=1}^n c_i x_i = 1$$

The last 3 are useful in showing expected value and variance properties:

$$E[\hat{\beta}_1] = \beta_1 \quad E[\hat{\beta}_0] = \beta_0$$

$$V[\hat{\beta}_1] = \frac{\sigma^2}{S_{xx}} \quad V[\hat{\beta}_0] = \sigma^2 \left(\frac{1}{n} + \frac{\bar{x}^2}{S_{xx}} \right)$$

The OLS Estimators are the **Best Linear Unbiased Estimators (BLUE)** by the **Gauss-Markov Theorem**, which states that:

In a linear regression model in which the errors

- have expectation zero,
- are uncorrelated, and
- have equal variances,

the best linear unbiased estimator (BLUE) of the coefficients is given by the ordinary least squares (OLS) estimator. Here, "best" means the estimator has the lowest variance as compared to other unbiased, linear estimators.

The errors do not need to be normal, nor do they need to be independent and identically distributed (only uncorrelated with mean zero and **homoscedastic** (i.e. all random variables have the same finite variance)). More useful properties of the least squares fit:

$$\sum_{i=1}^n y_i = \sum_{i=1}^n \hat{y}_i, \quad \sum_{i=1}^n (y_i - \hat{y}_i) = \sum_{i=1}^n \epsilon_i = 0, \quad \sum_{i=1}^n \epsilon_i x_i = \sum_{i=1}^n \epsilon_i \hat{y}_i = 0$$

The regression line also always passes through the centroid (\bar{y}, \bar{x}) of the data.

Estimation of σ^2 (2.2.3)

The **residual (error) sum of squares** (SS_{res}), is defined to be:

$$SS_{res} = \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \sum_{i=1}^n \epsilon_i^2$$

The **total sum of squares** is defined to be

$$SS_{total} = SS_{model} + SS_{res}$$

To estimate σ^2 , we use

$$\hat{\sigma}^2 = \frac{SS_{res}}{n-2} = MS_{res}$$

The quantity $n-2$ is the number of **degrees of freedom** (df) for the residual sum of squares. The df = $n-2$ because... ***

Since this estimate depends on the model and SS_{res} , any model error assumption violations could impact this estimate as well.

Hypothesis Testing on the Slope and Intercept

Three assumptions needed to apply procedures such as hypothesis testing and confidence intervals. Model errors, ϵ_i , are

- normally distributed,
- independently distributed, and
- have constant variance

i.e. $\epsilon_i \sim N(0, \sigma^2)$

Let's say we want to test if the slope ($\hat{\beta}_1$) is **NOT** equal to some constant, c .

This means we'd want to disprove the null hypothesis, H_0 , that $\hat{\beta}_1 = c$.

At this point, we'd need to calculate the **standard error** (aka standard deviation aka $\sqrt{V[\sigma^2]}$) of $\hat{\beta}_1$. This is defined like so:

$$se(\hat{\beta}_1) = \sqrt{\frac{MS_{res}}{S_{xx}}}$$

Our test statistic will then be:

$$t_0 = \frac{\hat{\beta}_1 - c}{se(\hat{\beta}_1)}$$

We reject H_0 (i.e. conclude there is sufficient evidence to believe that H_a is true) if:

$$|t_0| > t_{\frac{\alpha}{2}, n-2}$$

We can also use the p-value approach here as well.

To test if the intercept ($\hat{\beta}_0$) is **NOT** equal to some constant, c , we would do the same procedure except use $\hat{\beta}_0$'s standard error:

$$se(\hat{\beta}_0) = \sqrt{MS_{res} \left(\frac{1}{n} + \frac{\bar{x}^2}{S_{xx}} \right)}$$

Testing the significance of the regression (2.3.2)

$H_0: \beta_1 = 0$, $H_a: \beta_1 \neq 0$

This tests the significance of regression; that is, is there a linear relationship between the response and the regressor?

Failing to reject H_0 , implies that there is no linear relationship between y and x .

There is also an **analysis of variance** (ANOVA) approach.

SS_T (or SS_{Total}) = SS_{Model} or $SS_{Regression}$ or R + $SS_{Residual}$, so:

$$SS_T = SS_R + SS_{Res}, \quad \text{where } SS_R = \hat{\beta}_1 S_{xy}$$

$$df_T = df_R + df_{Res} \longrightarrow n-1 = 1 + (n-2)$$

Mean Squares:

$$MS_R = \frac{SS_R}{1}$$

$$MS_{Res} = \frac{SS_{Res}}{n-2}$$

10/8/17

for

$$y = x\beta + \epsilon,$$

$$\epsilon \sim N(0, \sigma^2 I)$$

Mid-term: Oct 19, Thursday Covers up to ch 4, some parts of ch 5 (depending)

- linear - constant variance - independent - Normal - $E(\epsilon) = 0$

4.4 Outliers:

- An outlier is an observation considerably different from the others - Formal tests for outliers - Points with large residuals may be outliers - Impact can be assessed by removing the points and refitting - How should they be treated?

4.5 Lack of Fit of the Regression Model

A formal test for lack of fit:

Assumes - Normality, independence, constant variance assumptions have been met - Only the first-order or straight line model is in doubt

Requires - replication of y for at least one level of x

With replication, we can obtain a "model-independent" estimate of σ^2

Say there are n_i observations of the response at the i th level of the regressor x_i , $i = 1, 2, \dots, m$

In other words, for a given x, you might have i occurrences of it.

y_{ij} denotes the j th observation on the response at x_i , $j = 1, 2, \dots, n_i$

Total number of observations is $n = \sum_{i=1}^m n_i$

i.e. if you have 3 occurrences of $x = 3.3$, and 3.3 is the n th x, then n_3 is 3.

$$SS_{total} (n - 1) = SS_{regr} (1) + SS_{resid} (n - 2)$$

$$SS_{resid} = SS_{pE} + SS_{LOF}$$

$$SS_{pE} = (y_{ij} - \bar{y}_i)^2$$

$$SS_{LOF} = (\bar{y}_i - \hat{y}_{ij})^2$$

$$SS_{res} = (y_{ij} - \hat{y}_{ij})^2$$

$$\sum_{i=1}^m \sum_{j=1}^{n_i} (y_{ij} - \hat{y}_{ij})^2 = \sum_{i=1}^m \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_i)^2 + \sum_{i=1}^m \sum_{j=1}^{n_i} (\bar{y}_i - \hat{y}_{ij})^2$$

$$(n - 2) = (n - m) + (m - 2)$$

$$SS_{res} = SS_{pE} + SS_{LOF}$$

How small is small?

Residual plots:

e_i vs \hat{y}_i

x_i vs \hat{e}_i

something vs normal

If you have 10 distinct data points, then m is 10. If you have 17 data points total (not necessarily distinct) then n is 17.

In R, you install the package `lme4`, and do `?pureErrorAnova`

10/12/17

Midterm will be Ch 0 - 5.

Ch 0 - Review (The Review HW and other basic stats. Specifically: estimating/hypothesizing/CI'ing μ and σ^2)

Ch 1 - Intro

Ch 2 - Simple Linear Regression

Ch 3 - Multiple Linear Regression

Ch 4 - Model Adequacy Checking / Diagnosis

Ch 5 - Transformation and Weighting to Correct Model Inadequacies

As far as Ch 5 goes

$$y = x\beta + \epsilon$$

Key assumptions:

Independence -

Constant Variance -

Normally distributed -

$$E[\epsilon] = 0$$

—

A linear regression model assumes: Linearity: $Y = \beta_0 + \beta_1 X$ Constant Variance: $\text{var}(Y|X) = \sigma^2$ Normality Dist. of Y s at any X is normal Independence Given X is, the Y is are independent

<http://www.columbia.edu/~so33/SusDev/Lecture2.pdf>

—

Linearity -

5.1 introduction

Data Transformation (fix the problem of assumption violation e.g. non-constant variance)

Generalized and Weighted Least Squares (fix the problem of assumption violation, e.g. correlated errors or non-constant variance)

Subject-Matter Knowledge (Ideally, the choice of metric should be made by the engineer or scientist with subject-matter knowledge)

5.2 Variance stabilizing Transformations

Constant variance assumption: - Often violated when the variance is **functionally related** to the mean i.e. $\sigma^2 \propto E[y]$ or $\sigma^2 = f(E[y])$

- Transformation on the response may eliminate the problem

-

—

An electric utility is interested in developing a model relating peak-hour demand (y) to total energy usage during the month (x).

This is an important planning problem because while most customers pay directly for energy usage (in kilowatt-hours), the generation system must be large enough to meet the maximum demand imposed.

Nonlinearity may be detected via the lack-of-fit test of Section 4.5

If a transformation of a nonlinear function can result in a linear function - we say it is intrinsically or transformably linear.

Ex:

$$y = \beta_0 e^{\beta_1 x} + \epsilon$$

$$\ln y = \ln \beta_0 + \beta_1 x + \ln \epsilon$$

—

10-17

More Ch 5 today.

For the midterm:

Basic stats (HW 2 stuff):

CI's, Hypothesis testing, Normal vs T for mean (μ).

Estimating σ^2 with χ^2 (CI, Hypoth testing)

Using s_1^2 and s_2^2 to estimate σ_1^2 and σ_2^2

calculate expectation, calculate variance

CH 4: Diagnosis

CH 5: Fix the problem

The key is Ch 3 and Ch 4

5 - know conceptually cuz no time to practice

So, essentially: Review, Ch 3, Ch 4

ACTUAL LECTURE:

$$y = x\beta + \epsilon$$

Assumptions:

linear: all the β 's power should be 1: β^1

Independence:

$$E[\epsilon] = 0$$

Normality

Constant variance

Residual values against regressors to see if they're correlated (checks to see if variance is constant)

If normal distribution is not satisfied (what does this mean?) then do transformations.

If you plot some y against some x and you see a curvy pattern,

and if you plot the residuals and you see a hump,

....

If our residuals have uncorrelated but not constant variance, then on the matrix e_1, e_1, e_3 etc on the columns, e_1, e_2, e_3 etc on the rows, the covariance of e_1, e_1 (aka the variance of e_1) will be, say, $2\sigma^2$, variance of e_2 is $4\sigma^2$, e_3 is $10\sigma^2$ etc, but the covariance of e_i, e_j where $i \neq j$ is still close to 0.

The weighted least square function multiplies the above matrix by a diagonal matrix containing the inverses of each of the diagonals from the above matrix

$\epsilon \sim N(0, \sigma^2 V)$ where V is some diagonal matrix containing the inverses above.

$E[\epsilon]$ is still 0

$V[\epsilon]$ will now be σ^2

$V[AX] = A V[X] A'$ if matrix, $\sigma^2 V[X]$ if constant.

$$\epsilon_1 = A\epsilon$$

$$V[\epsilon_1] = A V[\epsilon] A' = A\sigma^2 V A' \rightarrow \sigma^2 I$$

Positive Definite Matrix

$$K'K = KK' = V$$

Because ϵ is a covariance matrix, \exists a "square root" matrix K such that $KK' = V = K'K$

We do such a transformation: $K^{-1}\epsilon$

$$V(K^{-1}\epsilon) = K^{-1}\sigma^2 V(K^{-1})'$$

$$= \sigma^2 K^{-1}KK(K^{-1})'$$

A inverse transpose = A Transpose Inverse if symmetric.

So,

$$K^{-1}Y = K^{-1}X\beta + K^{-1}\epsilon$$

where

$$V[\epsilon] = \sigma^2 V$$

$$V[K^{-1}\epsilon] = \sigma^2 I$$

The least squares normal equations are:

... missed slide

$$\beta = (X'X)^{-1}X'Y$$

If weighted least squares is not implemented, but should have been:

BLUE: best linear unbiased estimator:

$$\hat{\beta} = (X'X)^{-1}X'y$$

$$E[\hat{\beta}] = \beta$$

The partial F-Test (extra sum of square)

If you have regressors x2, x7, x8

you have the full model: $\text{lm}(y \sim x2 + x7 + x8)$

$$SS_{\text{total}} = SS_{\text{reg}}(\text{full}) + SS_{\text{res}}(\text{Full})$$

Reduced Model:

$$SS_{\text{Total}}(\text{same as full model since its total variation in your data}) = SS_{\text{reg}}(\text{reduced}) + SS_{\text{res}}(\text{Reduced})$$

$$\text{lm}(y \sim x2 + x8)$$

Your full model's residuals should be smaller than your reduced model's residuals since your full model has more variables (aka R^2 is larger)

The difference between how many residuals the full has and how many residuals the reduced one has is the amount of residuals x7 is responsible for

$$F \text{ statistic} = (SS_{\text{reg}}(\text{full}) - SS_{\text{reg}}(\text{reduced})) \text{ divided by } MS_{\text{res}}(n - p) = F1, n - p$$

$$SS_{\text{res}}(\text{reduced}) - SS_{\text{res}}(\text{Full}) \text{ is the same as } (SS_{\text{reg}}(\text{full}) - SS_{\text{reg}}(\text{reduced}))$$

Partial regression plot:

$$y \sim x1, x2, x3, x4$$

$$y \sim x3 \text{ (you can just draw a scatter plot, that is one way)}$$

The partial regression plot:

$$\text{lm}(y \sim x1 + x2 + x4) - \text{get residual (1)} \quad \text{lm}(x3 \sim x1 + x2 + x4) - \text{get residual (2)}$$

then plot:

residual 1 vs residual 2

We do this because there might be correlations between variables, or there are some variations in x3 that could be explained by other variables.

We do this to account for those factors.

Partial F test vs partial regression plot? They both study relationships between y and other variables.

if residual 1 vs residual 2 is a linear pattern then y and x3 have a linear relationship and it should be included

if the plot is just noise (if there's a band) then there is no relationship and you probably shouldn't include it

4.2.2 - 4.2.4 4.2.6 4.4 4.5.1

5.1 5.2 5.3 5.5:1, 2, 3, 5